

# Intermediate Programming

## Day 11

# Outline

- Dynamic memory allocation
- Valgrind
- Review questions

# Returning an array in C

Q: Why doesn't this code work?

```
#include <stdio.h>
int *getArray3( void )
{
    int a[3];
    for( int i=0 ; i<3 ; i++ ) a[i] = 1;
    return a;
}
int main( void )
{
    int *list = NULL;
    list = getArray3();
    for( int i=0 ; i<3 ; i++ )
        printf( "%d " , list[i] );
    printf( "\n" );
    return 0;
}
```

```
>> ./a.out
Segmentation fault (core dumped)
>>
```



memory

# Stack frame


- Each function call gets a new, fixed size, *stack frame* “pushed onto” the stack in memory for storing
  - Local variables
  - Copies of function parameters
  - Information about the calling function
  - Other miscellany
- Once the current function call returns, the stack frame is “popped off” the stack, and execution returns to the calling function (so the current function’s frame is always on top)

# Returning an array in C

Q: Why doesn't this code work?

- **Frames get pushed on the stack when the function is called**
- And are popped off when it returns

```
#include <stdio.h>
int *getArray3( void )
{
    int a[3];
    for( int i=0 ; i<3 ; i++ ) a[i] = 1;
    return a;
}
int main( void )
{
    int *list = NULL;
    list = getArray3();
    for( int i=0 ; i<3 ; i++ )
        printf( "%d " , list[i] );
    printf( "\n" );
    return 0;
}
```




# Returning an array in C

Q: Why doesn't this code work?

- **Frames get pushed on the stack when the function is called**
- And are popped off when it returns

```
#include <stdio.h>
int *getArray3( void )
{
    int a[3];
    for( int i=0 ; i<3 ; i++ ) a[i] = 1;
    return a;
}
int main( void )
{
    int *list = NULL;
    list = getArray3();
    for( int i=0 ; i<3 ; i++ )
        printf( "%d " , list[i] );
    printf( "\n" );
    return 0;
}
```




# Returning an array in C

Q: Why doesn't this code work?

- Frames get pushed on the stack when the function is called
- And are popped off when it returns

```
#include <stdio.h>
int *getArray3( void )
{
    int a[3];
    for( int i=0 ; i<3 ; i++ ) a[i] = 1;
    return a;
}
int main( void )
{
    int *list = NULL;
    list = getArray3();
    for( int i=0 ; i<3 ; i++ )
        printf( "%d " , list[i] );
    printf( "\n" );
    return 0;
}
```




# Returning an array in C

Q: Why doesn't this code work?

- Frames get pushed on the stack when the function is called
- And are popped off when it returns

```
#include <stdio.h>
int *getArray3( void )
{
    int a[3];
    for( int i=0 ; i<3 ; i++ ) a[i] = 1;
    return a;
}
int main( void )
{
    int *list = NULL;
    list = getArray3();
    for( int i=0 ; i<3 ; i++ )
        printf( "%d " , list[i] );
    printf( "\n" );
    return 0;
}
```






# Returning an array in C

Q: Why doesn't this code work?

- Frames get pushed on the stack when the function is called
- **And are popped off when it returns**

A: By the time `list` is assigned the address `a`, the stack frame holding the array values is popped off and the address is no longer valid.

```
#include <stdio.h>
int *getArray3( void )
{
    int a[3];
    for( int i=0 ; i<3 ; i++ ) a[i] = 1;
    return a;
}
int main( void )
{
    int *list = NULL;
    list = getArray3();
    for( int i=0 ; i<3 ; i++ )
        printf( "%d " , list[i] );
    printf( "\n" );
    return 0;
}
```



# Limitations of arrays allocated in a stack frame

- Arrays (“statically”) allocated on the stack frame have limitations
  - Arrays created on a called function’s frame are not accessible to the calling function once the function returns
    - Returning `a` is meaningless

```
#include <stdio.h>
int *getArray3( void )
{
    int a[3];
    for( int i=0 ; i<3 ; i++ ) a[i] = 1;
    return a;
}
int main( void )
{
    int *list = NULL;
    list = getArray3();
    for( int i=0 ; i<3 ; i++ )
        printf( "%d " , list[i] );
    printf( "\n" );
    return 0;
}
```

# Limitations of arrays allocated in a stack frame

- Arrays (“statically”) allocated on the stack frame have limitations
  - Arrays created on a called function’s frame are not accessible to the calling function once the function returns
  - Stack frames have fixed (small) size
    - Cannot statically allocate more memory than the size of the stack frame

```
#include <stdio.h>
int main( void )
{
    int values[100000000];
    return 0;
}
```

```
>> ./a.out
Segmentation fault (core dumped)
>>
```

# Dynamic Memory Allocation

- Dynamically-allocated memory:
  - Is located on the *heap* – a part of memory separate from the stack
  - Lives as long as we like (until the entire program ends)
    - We don't lose access to it when function call returns
      - ✓ We can return it to a calling function!
      - ✗ We are responsible for managing the memory (and cleaning it up when we're done)
  - Is not subject to frame stack's size limitations because it isn't part of the stack

# Dynamic Memory Allocation

- ✗ We are responsible for managing the memory
  - Memory should be *deallocated* when it is no longer needed
  - Allocated memory is not available to other programs/users until we deallocate it
  - Failing to deallocate memory is the cause of “memory leaks”

# Dynamic Memory Allocation (single)

- We allocate memory using the `malloc` command (in `stdlib.h`)

`void *malloc( size_t );`

- Input: how much memory (in bytes) to allocate
  - `size_t` is an unsigned integer type:
    - 4 bytes on 32-bit machines
    - 8 bytes on 64-bit machines
- Output: the location (on the heap) of the memory
  - `malloc` doesn't need to know what you're going to store just how much memory you need (so it returns a `void*`)

```
...  
int *ip = malloc( sizeof( int ) );
```

\*On a 32-bit machine you can't ask for more than  $2^{32}$  bytes.

# Dynamic Memory Allocation (single)

- We allocate memory using the `malloc` command (in `stdlib.h`)

`void *malloc( size_t );`

- Check that allocation succeeded

```
...
int *ip = malloc( sizeof( int ) );
if( ip==NULL )
{
    fprintf( stderr , "...\\n" );
    // do something
}
```

# Dynamic Memory Allocation (single)

- We allocate memory using the `malloc` command (in `stdlib.h`)

`void *malloc( size_t );`

- Check that allocation succeeded
- After allocation with `malloc`, memory cannot be assumed to be initialized

```
...
int *ip = malloc( sizeof( int ) );
if( ip==NULL )
{
    fprintf( stderr , "...\\n" );
    // do something
}
*ip = 0;
...
```



# Dynamic Memory Allocation (single)

- We allocate memory using the `malloc` command (in `stdlib.h`)

`void *malloc( size_t );`

- Check that allocation succeeded
- After allocation with `malloc`, memory cannot be assumed to be initialized
- When done using dynamically-allocated memory, deallocate using `free`

`void free( void *ptr );`

- Input: address of the memory on the heap

```
...
int *ip = malloc( sizeof( int ) );
if( ip==NULL )
{
    fprintf( stderr , "...\\n" );
    // do something
}
*ip = 0;
...
free( ip );
```

# Dynamic Memory Allocation (single)

- We allocate memory using the `malloc` command (in `stdlib.h`)

`void *malloc( size_t );`

- Check that allocation succeeded
- After allocation with `malloc`, memory cannot be assumed to be initialized
- When done using dynamically-allocated memory, deallocate using `free`

`void free( void *ptr );`

- It's good practice to set the pointer to `NULL` to avoid accidental use

```
...
int *ip = malloc( sizeof( int ) );
if( ip==NULL )
{
    fprintf( stderr , "...\\n" );
    // do something
}
*ip = 0;
...
if( ip!=NULL ) free( ip );
ip = NULL;
...
```

# Dynamic Memory Allocation (multiple)

- We allocate memory using the `malloc` command (in `stdlib.h`)

`void *malloc( size_t );`

- Check that allocation succeeded
- After allocation with `malloc`, memory cannot be assumed to be initialized
- When done using dynamically-allocated memory, deallocate using `free`

`void free( void *ptr );`

- It's good practice to set the pointer to `NULL` to avoid accidental use

```
...
int *ip = malloc( sizeof( int ) * sz );
if( ip==NULL )
{
    fprintf( stderr , "...\\n" );
    // do something
}
for( int i=0 ; i<sz ; i++ ) ip[i] = 0;
...
if( ip!=NULL ) free( ip );
ip = NULL;
...
```

# Dynamic Memory Allocation (multiple)

- We allocate memory using the `malloc` command (in `stdlib.h`)

`void *malloc( size_t );`

- Check that allocation succeeded
- After allocation with `malloc`, memory cannot be assumed to be initialized
- When done using dynamically-allocated memory, deallocate using `free`

`void free( void *ptr );`

Note:

Call `free` once for every `malloc` called.

```
...
int *ip = malloc( sizeof( int ) * sz );
if( ip==NULL )
{
    fprintf( stderr , "...\\n" );
    // do something
}
for( int i=0 ; i<sz ; i++ ) ip[i] = 0;
...
if( ip!=NULL ) free( ip );
ip = NULL;
...
```

# Deallocation

- Deallocation does not have to happen in the same function where allocation occurred. . .
  - But it does have to happen!
  - Otherwise you can get a *memory leak*

[BADNESS] Only the last allocation was deallocated!


```
#include <stdio.h>
#include <stdlib.h>
int *getArray( unsigned int sz )
{
    int *a = malloc( sizeof(int) * sz );
    for( int i=0 ; i<sz ; i++ ) a[i] = 1;
    return a;
}
int main( void )
{
    int *list;
    for( int i=0 ; i<1000000 ; i++ )
    {
        list = getArray( 100000 );
        // do something with list
    }
    free( list );
    return 0;
}
```

# Deallocation

- Deallocation does not have to happen in the same function where allocation occurred. . .
  - But it does have to happen!
  - Otherwise you can get a *memory leak*

```
#include <stdio.h>
#include <stdlib.h>
int *getArray( unsigned int sz )
{
    int *a = malloc( sizeof(int) * sz );
    for( int i=0 ; i<sz ; i++ ) a[i] = 1;
    return a;
}
int main( void )
{
    int *list;
    for( int i=0 ; i<1000000 ; i++ )
    {
        list = getArray( 100000 );
        // do something with list
        free( list );
    }
    return 0;
}
```

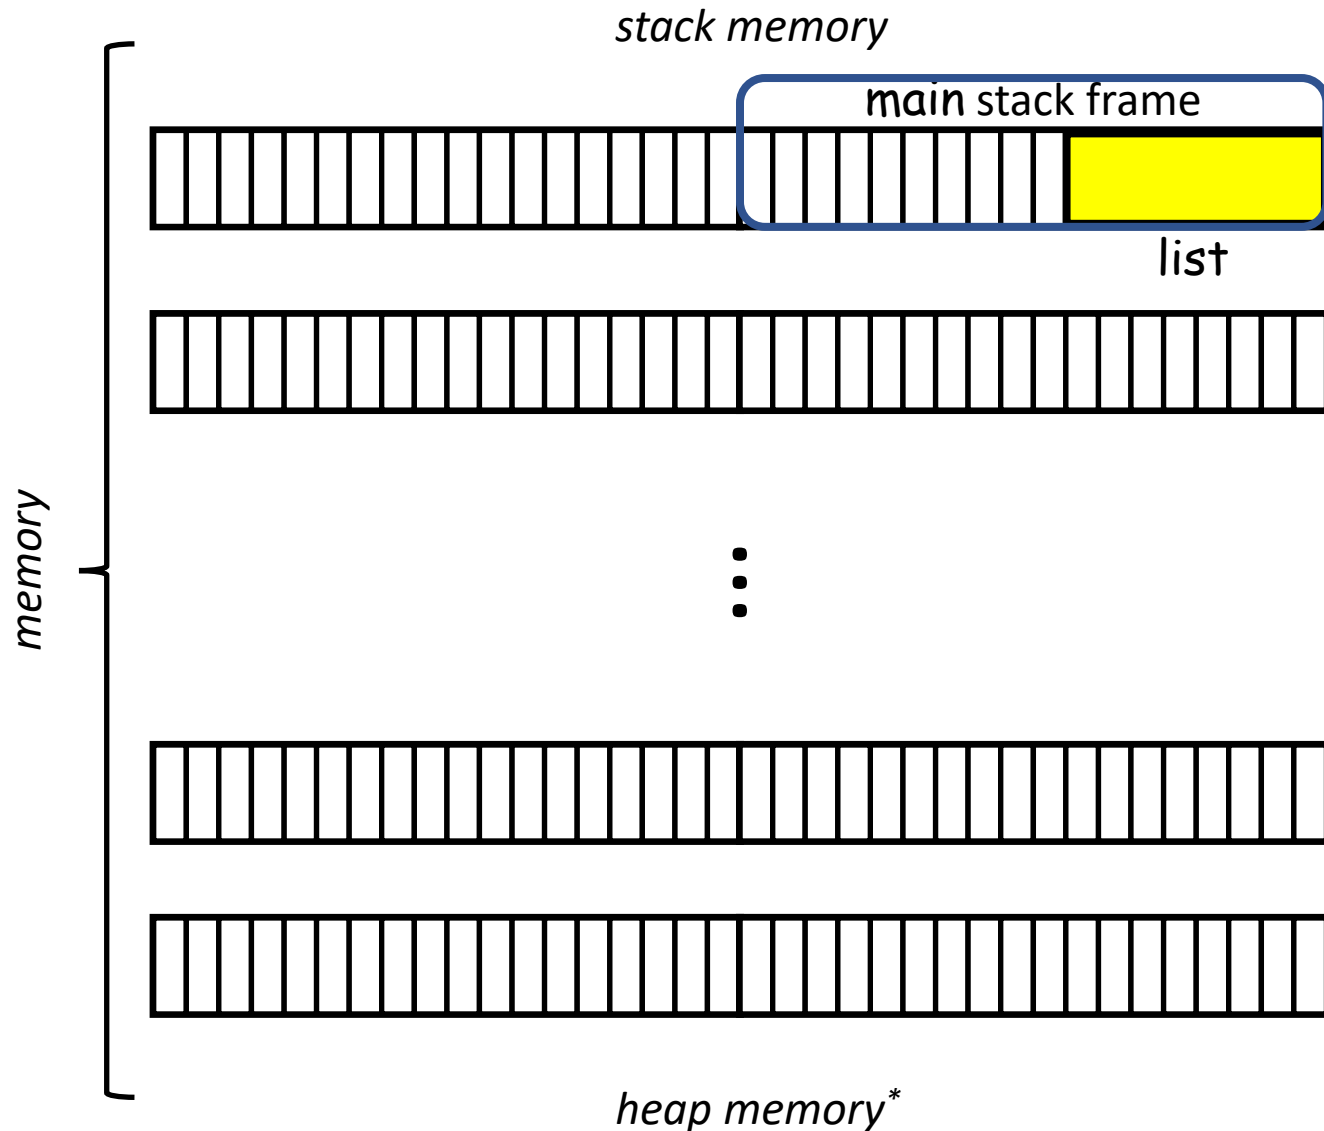
# Dynamic Memory Allocation



```
#include <stdio.h>
#include <stdlib.h>
int *getArray( unsigned int sz )
{
    int *a;
    a = malloc( sizeof(int) * sz );
    for( int i=0 ; i<sz ; i++ ) a[i] = 0;
    return a;
}
int main( void )
{
    int *list;
    list = getArray( 8 );
    // do something with list
    free( list );
    list = NULL;
    return 0;
}
```

\* Visualization is not to scale:  
There is much more room on the heap than on the stack.

# Dynamic Memory Allocation



```
#include <stdio.h>
#include <stdlib.h>

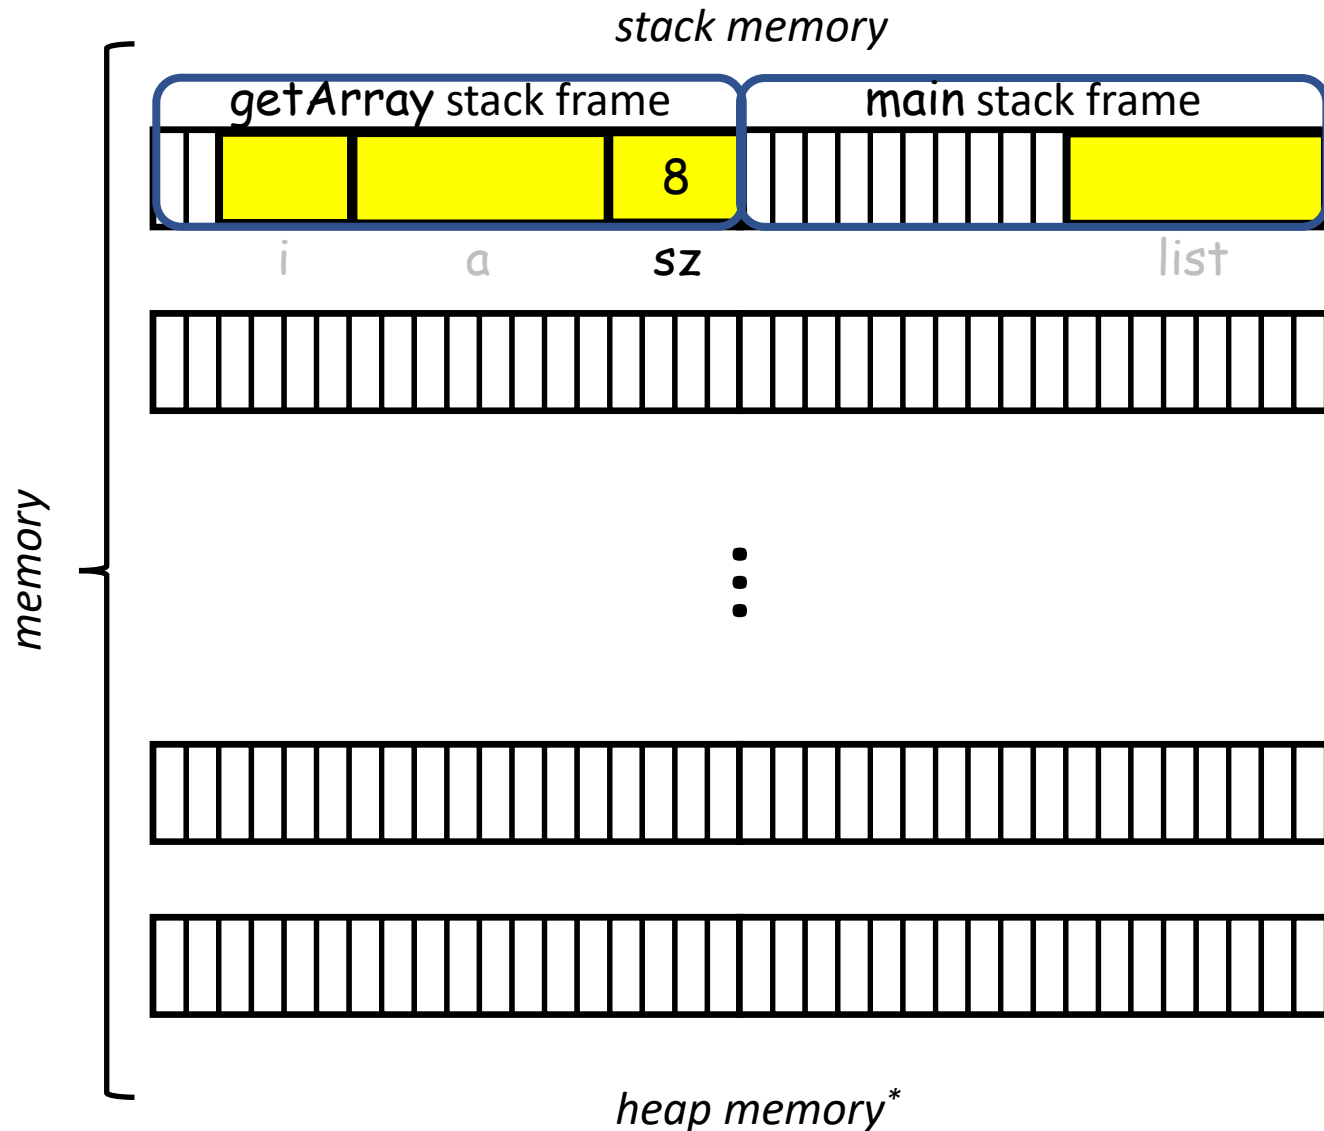
int *getArray( unsigned int sz )
{
    int *a;
    a = malloc( sizeof(int) * sz );
    for( int i=0 ; i<sz ; i++ ) a[i] = 0;
    return a;
}

int main( void )
{
    int *list;
    list = getArray( 8 );
    // do something with list
    free( list );
    list = NULL;
    return 0;
}
```

\* Visualization is not to scale:  
There is much more room on the heap than on the stack.



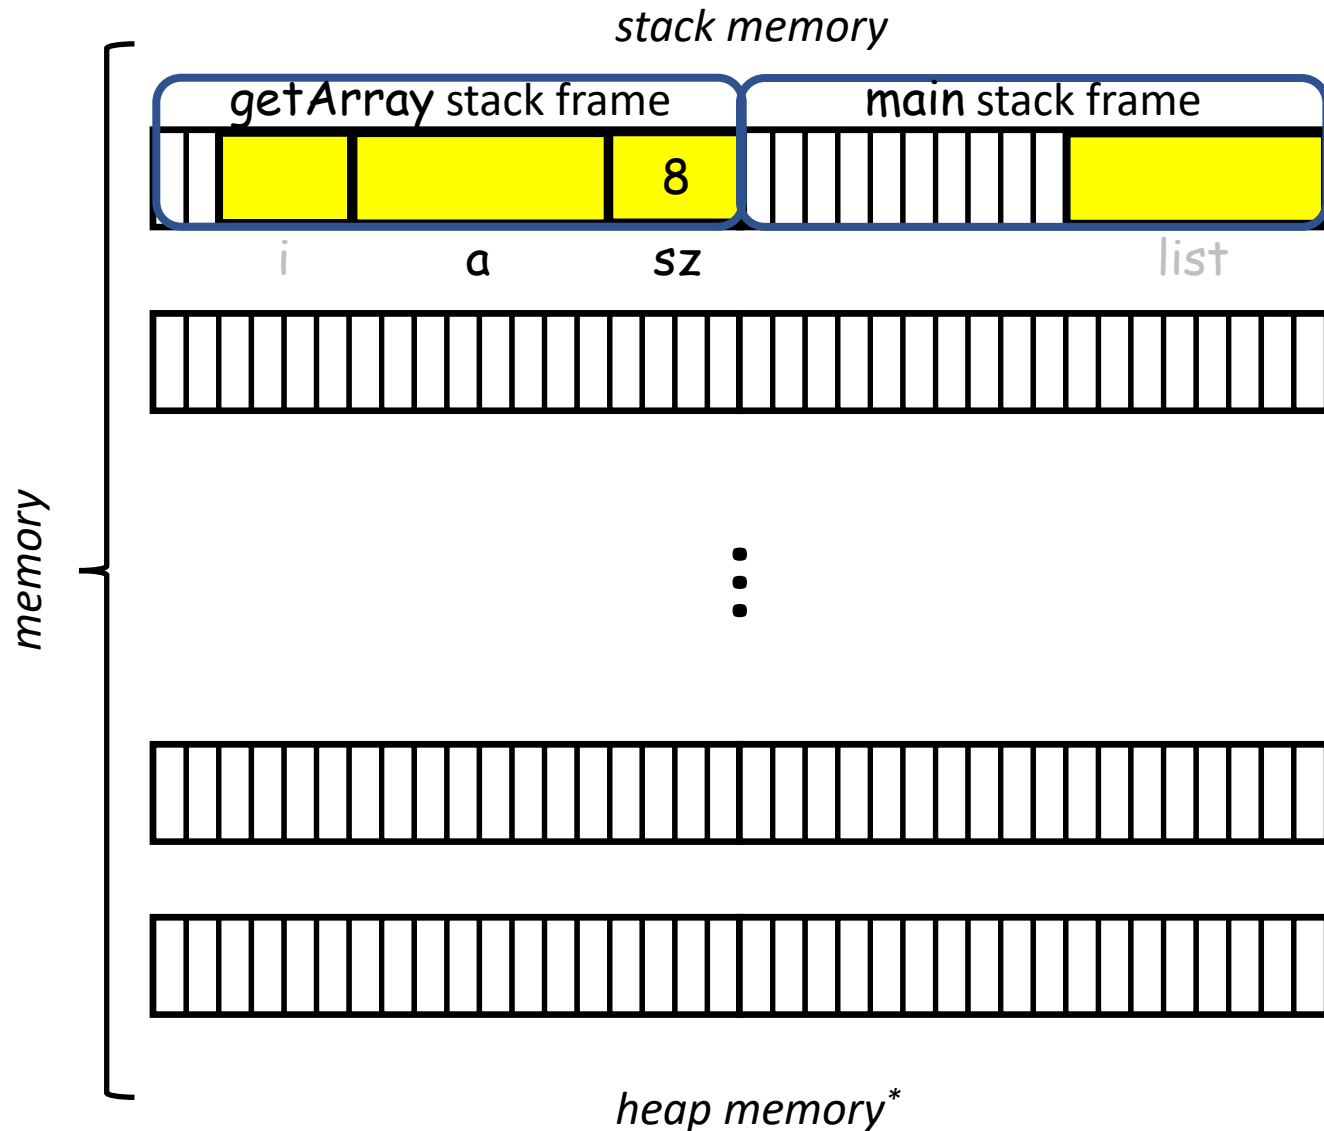
# Dynamic Memory Allocation



```
#include <stdio.h>
#include <stdlib.h>
int *getArray( unsigned int sz )
{
    int *a;
    a = malloc( sizeof(int) * sz );
    for( int i=0 ; i<sz ; i++ ) a[i] = 0;
    return a;
}
int main( void )
{
    int *list;
    list = getArray( 8 );
    // do something with list
    free( list );
    list = NULL;
    return 0;
}
```

\* Visualization is not to scale:  
There is much more room on the heap than on the stack.

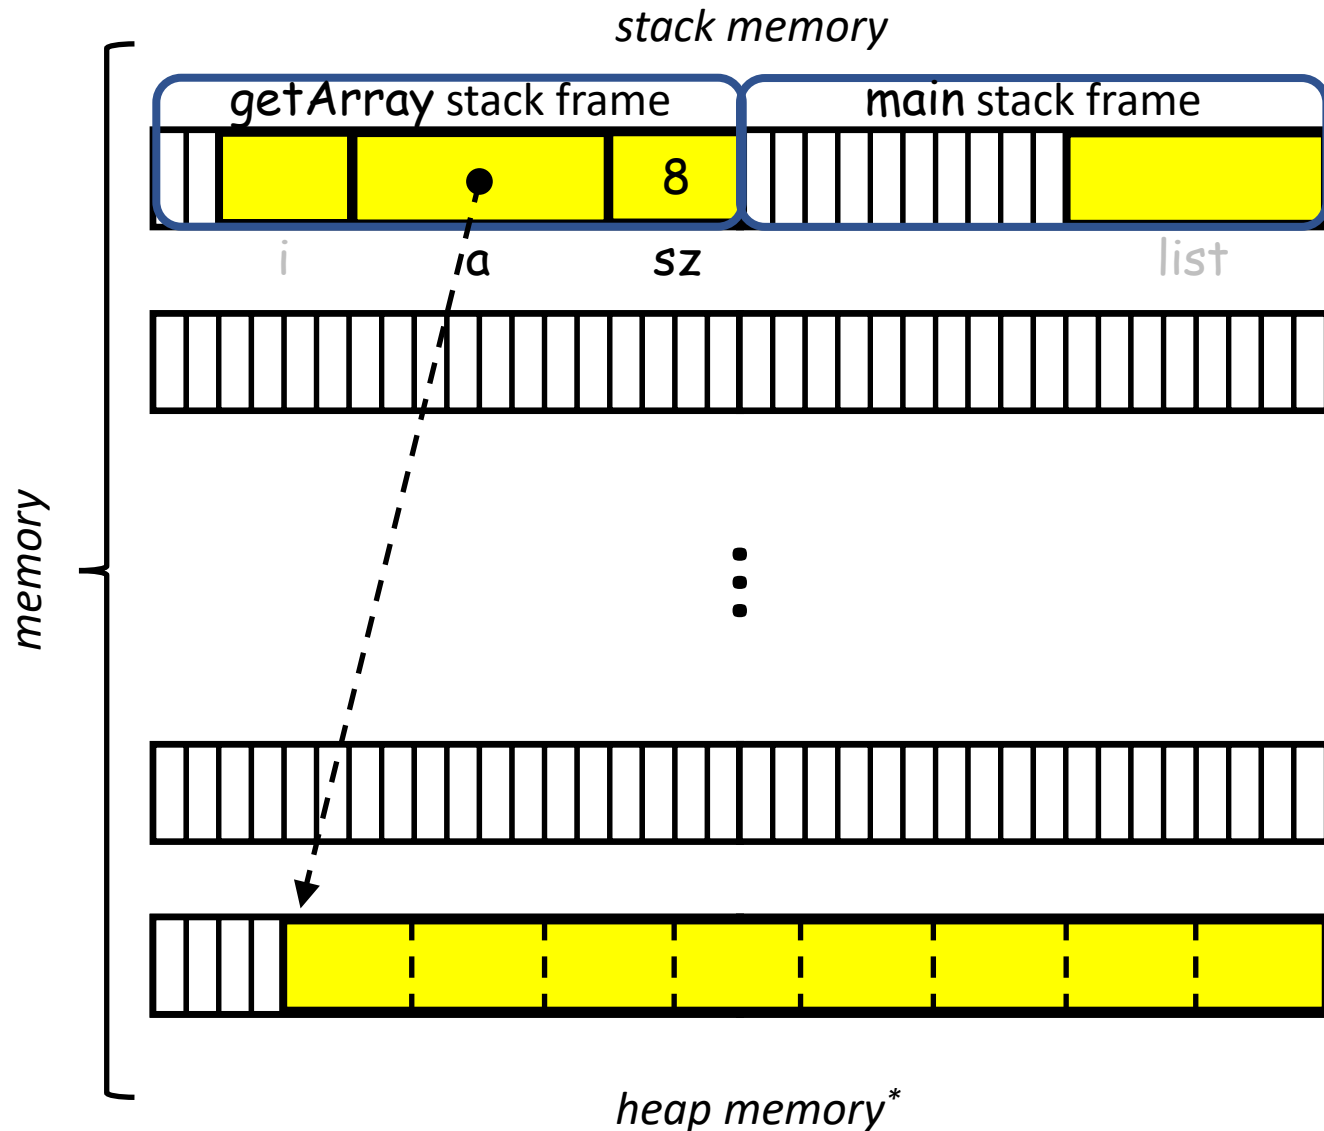
# Dynamic Memory Allocation



```
#include <stdio.h>
#include <stdlib.h>
int *getArray( unsigned int sz )
{
    int *a;
    a = malloc( sizeof(int) * sz );
    for( int i=0 ; i<sz ; i++ ) a[i] = 0;
    return a;
}
int main( void )
{
    int *list;
    list = getArray( 8 );
    // do something with list
    free( list );
    list = NULL;
    return 0;
}
```

\* Visualization is not to scale:  
There is much more room on the heap than on the stack.

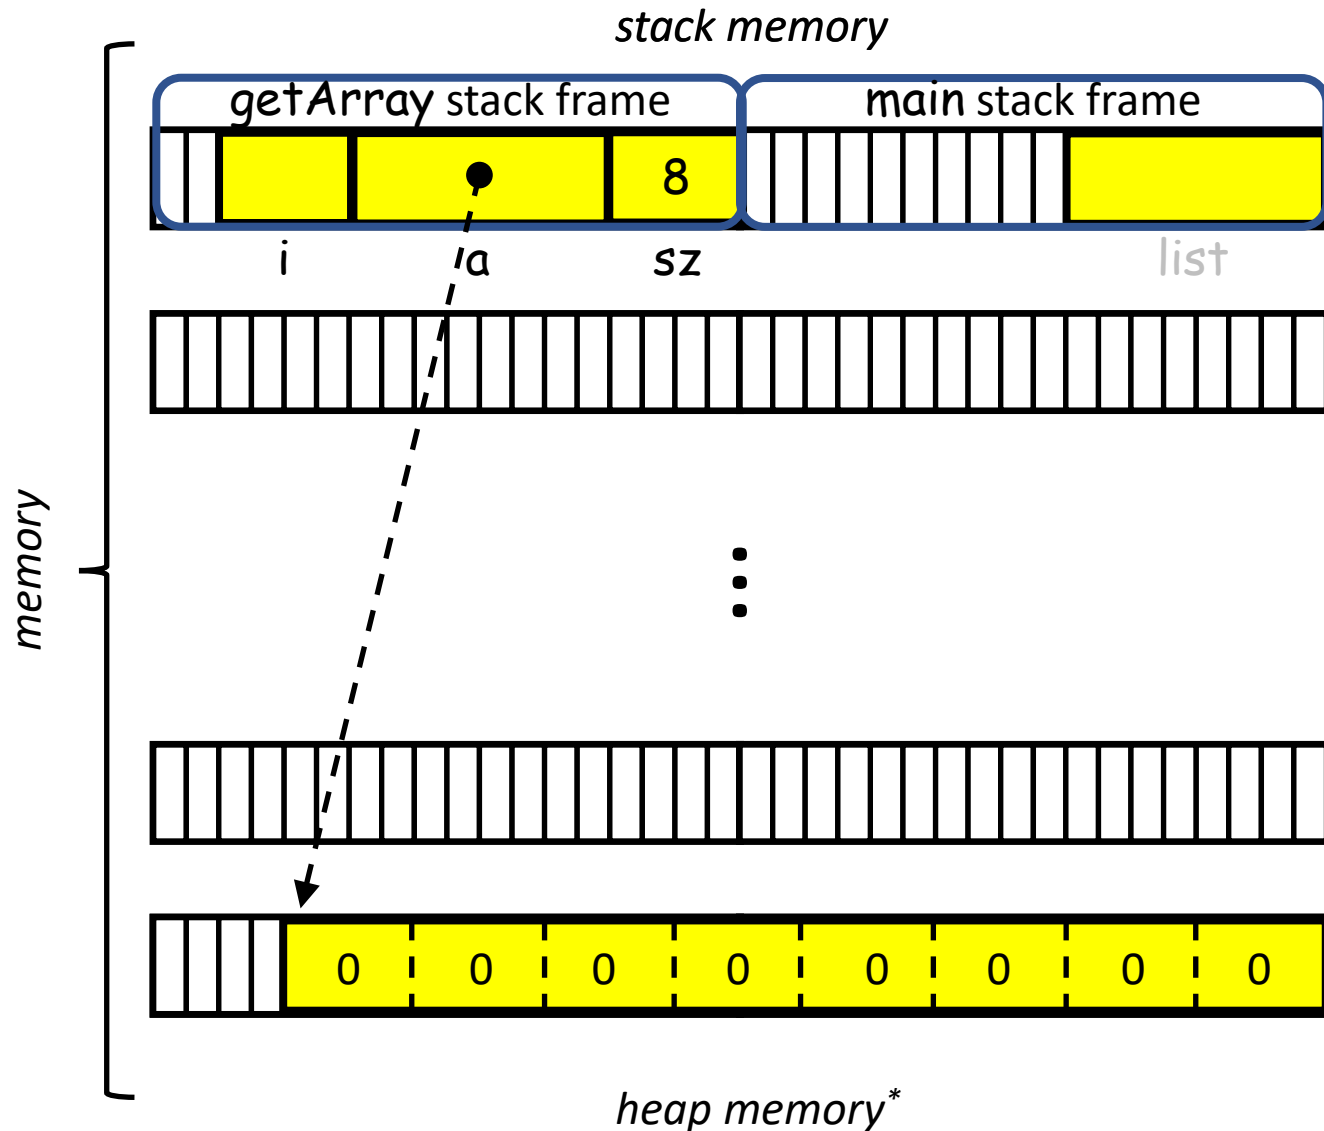
# Dynamic Memory Allocation



```
#include <stdio.h>
#include <stdlib.h>
int *getArray( unsigned int sz )
{
    int *a;
    a = malloc( sizeof(int) * sz );
    for( int i=0 ; i<sz ; i++ ) a[i] = 0;
    return a;
}
int main( void )
{
    int *list;
    list = getArray( 8 );
    // do something with list
    free( list );
    list = NULL;
    return 0;
}
```

\* Visualization is not to scale:  
There is much more room on the heap than on the stack.

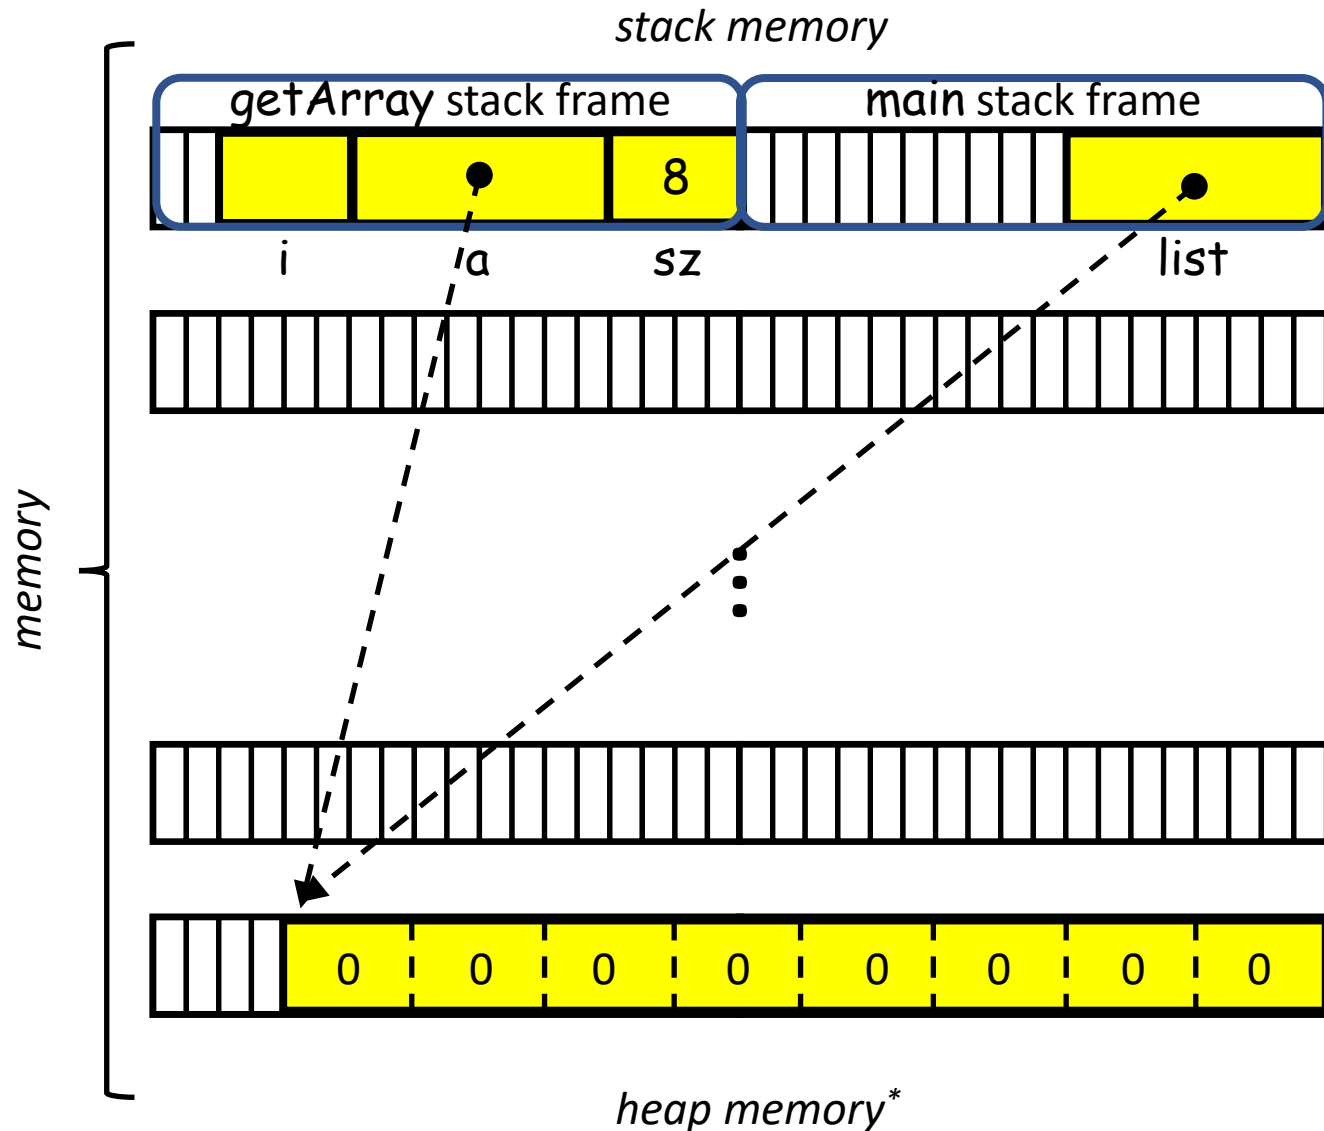
# Dynamic Memory Allocation



```
#include <stdio.h>
#include <stdlib.h>
int *getArray( unsigned int sz )
{
    int *a;
    a = malloc( sizeof(int) * sz );
    for( int i=0 ; i<sz ; i++ ) a[i] = 0;
    return a;
}
int main( void )
{
    int *list;
    list = getArray( 8 );
    // do something with list
    free( list );
    list = NULL;
    return 0;
}
```

\* Visualization is not to scale:  
There is much more room on the heap than on the stack.


# Dynamic Memory Allocation



```
#include <stdio.h>
#include <stdlib.h>
int *getArray( unsigned int sz )
{
    int *a;
    a = malloc( sizeof(int) * sz );
    for( int i=0 ; i<sz ; i++ ) a[i] = 0;
    return a;
}
int main( void )
{
    int *list;
    list = getArray( 8 );
    // do something with list
    free( list );
    list = NULL;
    return 0;
}
```

\* Visualization is not to scale:  
There is much more room on the heap than on the stack.

# Dynamic Memory Allocation




```
#include <stdio.h>
#include <stdlib.h>

int *getArray( unsigned int sz )
{
    int *a;
    a = malloc( sizeof(int) * sz );
    for( int i=0 ; i<sz ; i++ ) a[i] = 0;
    return a;
}

int main( void )
{
    int *list;
    list = getArray( 8 );
    // do something with list
    free( list );
    list = NULL;
    return 0;
}
```

\* Visualization is not to scale:  
There is much more room on the heap than on the stack.

# Dynamic Memory Allocation




```
#include <stdio.h>
#include <stdlib.h>

int *getArray( unsigned int sz )
{
    int *a;
    a = malloc( sizeof(int) * sz );
    for( int i=0 ; i<sz ; i++ ) a[i] = 0;
    return a;
}

int main( void )
{
    int *list;
    list = getArray( 8 );
    // do something with list
    free( list );
    list = NULL;
    return 0;
}
```

\* Visualization is not to scale:  
There is much more room on the heap than on the stack.

# Dynamic Memory Allocation



```
#include <stdio.h>
#include <stdlib.h>


int *getArray( unsigned int sz )
{
    int *a;
    a = malloc( sizeof(int) * sz );
    for( int i=0 ; i<sz ; i++ ) a[i] = 0;
    return a;
}

int main( void )
{
    int *list;
    list = getArray( 8 );
    // do something with list
    free( list );
    list = NULL;
    return 0;
}
```

\* Visualization is not to scale:  
There is much more room on the heap than on the stack.



# Dynamic Memory Allocation




```
#include <stdio.h>
#include <stdlib.h>

int *getArray( unsigned int sz )
{
    int *a;
    a = malloc( sizeof(int) * sz );
    for( int i=0 ; i<sz ; i++ ) a[i] = 0;
    return a;
}

int main( void )
{
    int *list;
    list = getArray( 8 );
    // do something with list
    free( list );
    list = NULL;
    return 0;
}
```

\* Visualization is not to scale:  
There is much more room on the heap than on the stack.

# Dynamic Memory Allocation



```
#include <stdio.h>
#include <stdlib.h>
int *getArray( unsigned int sz )
{
    int *a;
    a = malloc( sizeof(int) * sz );
    for( int i=0 ; i<sz ; i++ ) a[i] = 0;
    return a;
}
int main( void )
{
    int *list;
    list = getArray( 8 );
    // do something with list
    free( list );
    list = NULL;
    return 0;
}
```

\* Visualization is not to scale:  
There is much more room on the heap than on the stack.

# Working w/ Memory (`stdlib.h`)

- Allocate and clear memory:

`void *calloc( size_t num , size_t size );`

Input:

- **num**: the number of elements
- **size**: the size of each element (undefined behavior if zero)

Output:

- On success: a pointer to the allocated and cleared memory
- On failure: a null-pointer

```
#include <stdio.h>
#include <stdlib.h>
int main( void )
{
    int *v = calloc( 4 , sizeof(int) );
    if( !v ) ...
    for( int i=0 ; i<4 ; i++ ) printf( "%d] %d\n" , i , v[i] );
    free( v );
    return 0;
}
```

```
>> ./a.out
0] 0
1] 0
2] 0
3] 0
>>
```

# Working w/ Memory (`stdlib.h`)

- Change the size of memory pointed to by a pointer while preserving the data (up to the minimum of the old and new sizes):

`void *realloc( void *ptr , size_t size );`

Input:

- `ptr`: pointer to memory on the heap (if `NULL`, `realloc` behaves like `malloc`)
- `size`: new size for the memory block (if 0, `realloc` behaves like `free`)

Output:

- On failure or if `size=0`: a null-pointer
- On success: a pointer to the new memory with old values preserved up to `num` bytes (may or may not be the address of the old memory)

```
#include <stdio.h>
#include <stdlib.h>
int main( void )
{
    int *v = malloc( sizeof(int)*3 );
    for( int i=0 ; i<3 ; i++ ) v[i] = i;
    v = realloc( v , sizeof(int)*4 );
    v[3] = 3;
    for( int i=0 ; i<4 ; i++ )
        printf( "%d] %d\n" , i , v[i] );
    free( v );
    return 0;
}
```

```
>> ./a.out
0] 0
1] 1
2] 2
3] 3
>>
```

# Working w/ Memory (`stdlib.h`)

- Change the size of memory pointed to by a pointer while preserving the data (up to the minimum of the old and new sizes):

`void *realloc( void *ptr , size_t size );`

## WARNING:

`realloc` will try to extend/shrink the memory in place but is not guaranteed to (e.g. if there isn't sufficient space to grow).

⇒ Do not assume that the input and output pointers are equal

- In that case, `realloc` handles deallocation of the old memory (e.g. there is no call "`free(v);`")

- On failure or if `size=0`: a null-pointer
- On success: a pointer to the new memory with old values preserved up to `num` bytes (may or may not be the address of the old memory)

```
#include <stdio.h>
#include <stdlib.h>
int main( void )
{
```

```
    int *v = malloc( sizeof(int)*3 );
    for( int i=0 ; i<3 ; i++ ) v[i] = i;
    v = realloc( v , sizeof(int)*4 );
    v[3] = 3;
    for( int i=0 ; i<4 ; i++ )
        printf( "%d] %d\n" , i , v[i] );
    free( v );
    return 0;
```

```
}
```

```
>> ./a.out
0] 0
1] 1
2] 2
3] 3
>>
```

# Outline

- Dynamic memory allocation
- **Valgrind**
- Review questions

# Valgrind

- Easy-to-use tool for finding memory leaks and other memory issues
- Compile with -g to get more helpful output from valgrind
- Then run using valgrind:

```
valgrind --leak-check=full ./myFile <arg1> <arg2> ...
```

```
#include <stdio.h>
int main( void )
{
    printf( "Hello world!\n" );
    return 0;
}
```

```
>> gcc -std=c99 -Wall -Wextra -g foo.c
>> ./a.out
Hello world!
>>
```

# Valgrind

- Easy-to-use tool for finding memory leaks and other memory issues
- Compile with -g to get more helpful output from valgrind
- Then run using valgrind:

```
va >> valgrind --leak-check=full ./a.out
```

*header*

```
==12133== Command: ./a.out
```

```
==12133==
```

```
Hello World!
```

```
==12133==
```

```
==12133== HEAP SUMMARY:
```

```
==12133==    in use at exit: 0 bytes in 0 blocks
```

```
==12133== total heap usage: 1 allocs, 1 frees, 1,024 bytes allocated
```

```
==12133==
```

```
==12133== All heap blocks were freed -- no leaks are possible
```

```
==12133==
```

```
==12133== For counts of detected and suppressed errors, rerun with: -v
```

```
==12133== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```

```
>>
```

>

```
world!\n" );
```

```
Wextra -g foo.c
```

See also <http://valgrind.org/>



# Valgrind

- Your program output is interspersed with messages from `valgrind`
- Kinds of issues flagged by `valgrind`
  - Invalid reads or writes:  
Attempts to dereference pointers to memory that's not yours
  - Memory leaks:  
Failing to deallocate a block of memory you allocate previously
    - Info about leaks appears in HEAP SUMMARY section

# Valgrind

```
1. #include <stdio.h>
2. #include <stdlib.h>
3. #include <string.h>
4. #include <assert.h>
5. char *string_copy( const char *in )
6. {
7.     char *out = malloc( strlen( in ) );
8.     assert( out!= NULL );
9.     return strcpy( out , in );
10. }
11. int main( void )
12. {
13.     char *str = string_copy( "hello" );
14.     assert( str!=NULL );
15.     printf( "%s\n" , str );
16.     return 0;
17. }
```

```
>> ./a.out
hello
>>
```

```
>> valgrind --leak-check=full ./a.out
```

*header*

```
==17647== Command: ./a.out
```

```
==17647==
```

```
==17647== Invalid write of size 1
```

```
==17647==    at 0x4C30CB7: strcpy (vg_replace_strmem.c:506)
```

```
==17647==    by 0x40067C: string_copy (foo.c:9)
```

```
==17647==    by 0x400690: main (foo.c:13)
```

```
==17647== Address 0x5200045 is 0 bytes after a block of size 5 alloc'd
```

```
==17647==    at 0x4C2DB9D: malloc (vg_replace_malloc.c:299)
```

```
==17647==    by 0x400645: string_copy (foo.c:7)
```

```
==17647==    by 0x400690: main (foo.c:13)
```

```
==17647==
```

```
==17647== Invalid read of size 1
```

```
==17647==    at 0x4C30BC4: strlen (vg_replace_strmem.c:454)
```

```
==17647==    by 0x4EAAA1: puts (in /usr/lib64/libc-2.24.so)
```

```
==17647==    by 0x4006C0: main (foo.c:15)
```

```
==17647== Address 0x5200045 is 0 bytes after a block of size 5 alloc'd
```

```
==17647==    at 0x4C2DB9D: malloc (vg_replace_malloc.c:299)
```

```
==17647==    by 0x400645: string_copy (foo.c:7)
```

```
==17647==    by 0x400690: main (foo.c:13)
```

```
==17647==
```

```
hello
```

```
==17647==
```

*first half of output*

# Valgrind

```
1. #include <stdio.h>
2. #include <stdlib.h>
3. #include <string.h>
4. #include <assert.h>
5. char *string_copy( const char *in )
6. {
7.     char *out = malloc( strlen( in ) );
8.     assert( out!= NULL );
9.     return strcpy( out , in );
10. }
11. int main( void )
12. {
13.     char *str = string_copy( "hello" );
14.     assert( str!=NULL );
15.     printf( "%s\n" , str );
16.     return 0;
17. }
```

```
>> ./a.out
hello
>>
```

*second half of output*

```
==17647== HEAP SUMMARY:
==17647==      in use at exit: 5 bytes in 1 blocks
==17647==    total heap usage: 2 allocs, 1 frees, 1,029 bytes allocated
==17647==
==17647== 5 bytes in 1 blocks are definitely lost in loss record 1 of 1
==17647==    at 0x4C2DB9D: malloc (vg_replace_malloc.c:299)
==17647==    by 0x400645: string_copy (foo.c:7)
==17647==    by 0x400690: main (foo.c:13)
==17647==
==17647== LEAK SUMMARY:
==17647==    definitely lost: 5 bytes in 1 blocks
==17647==    indirectly lost: 0 bytes in 0 blocks
==17647==    possibly lost: 0 bytes in 0 blocks
==17647==    still reachable: 0 bytes in 0 blocks
==17647==          suppressed: 0 bytes in 0 blocks
==17647==
==17647== For counts of detected and suppressed errors, rerun with: -v
==17647== ERROR SUMMARY: 3 errors from 3 contexts (suppressed: 0 from 0)
>>
```

# Valgrind

```
1. #include <stdio.h>
2. #include <stdlib.h>
3. #include <string.h>
4. #include <assert.h>
5. char *string_copy( const char *in )
6. {
7.     char *out = malloc( strlen( in ) );
8.     assert( out!= NULL );
9.     return strcpy( out , in );
10. }
11. int main( void )
12. {
13.     char *str = string_copy( "hello" );
14.     assert( str!=NULL );
15.     printf( "%s\n" , str );
16.     return 0;
17. }
```

```
>> ./a.out
hello
>>
```

```
>> valgrind --leak-check=full ./a.out
```

*header*

```
==17647== Command: ./a.out
```

```
==17647==
```

```
==17647== Invalid write of size 1
```

```
==17647==    at 0x4C30CB7: strcpy (vg_replace_strmem.c:506)
```

```
==17647==    by 0x40067C: string_copy (foo.c:9)
```

```
==17647==    by 0x400690: main (foo.c:13)
```

```
==17647== Address 0x5200045 is 0 bytes after a block of size 5 alloc'd
```

```
==17647==    at 0x4C2DB9D: malloc (vg_replace_malloc.c:299)
```

```
==17647==    by 0x400645: string_copy (foo.c:7)
```

```
==17647==    by 0x400690: main (foo.c:13)
```

```
==17647==
```

```
==17647== Invalid read of size 1
```

```
==17647==    at 0x4C30BC4: strlen (vg_replace_strmem.c:454)
```

```
==17647==    by 0x4EAAA1: puts (in /usr/lib64/libc-2.24.so)
```

```
==17647==    by 0x4006C0: main (foo.c:15)
```

```
==17647== Address 0x5200045 is 0 bytes after a block of size 5 alloc'd
```

```
==17647==    at 0x4C2DB9D: malloc (vg_replace_malloc.c:299)
```

```
==17647==    by 0x400645: string_copy (foo.c:7)
```

```
==17647==    by 0x400690: main (foo.c:13)
```

```
==17647==
```

```
hello
```

```
==17647==
```

*first half of output*

# Valgrind

```
1. #include <stdio.h>
2. #include <stdlib.h>
3. #include <string.h>
4. #include <assert.h>
5. char *string_copy( const char *in )
6. {
7.     char *out = malloc( strlen( in ) );
8.     assert( out!= NULL );
9.     return strcpy( out , in );
10. }
11. int main( void )
12. {
13.     char *str = string_copy( "hello" );
14.     assert( str!=NULL );
15.     printf( "%s\n" , str );
16.     return 0;
17. }
```

```
>> ./a.out
hello
>>
```

```
>> valgrind --leak-check=full ./a.out
```

*header*

```
==17647== Command: ./a.out
```

```
==17647==
```

```
==17647== Invalid write of size 1
```

```
==17647==    at 0x4C30CB7: strcpy (vg_replace_strmem.c:506)
```

```
==17647==    by 0x40067C: string_copy (foo.c:9)
```

```
==17647==    by 0x400690: main (foo.c:13)
```

```
==17647== Address 0x5200045 is 0 bytes after a block of size 5 alloc'd
```

```
==17647==    at 0x4C2DB9D: malloc (vg_replace_malloc.c:299)
```

```
==17647==    by 0x400645: string_copy (foo.c:7)
```

```
==17647==    by 0x400690: main (foo.c:13)
```

```
==17647==
```

```
==17647== Invalid read of size 1
```

```
==17647==    at 0x4C30BC4: strlen (vg_replace_strmem.c:454)
```

```
==17647==    by 0x4EAAA1: puts (in /usr/lib64/libc-2.24.so)
```

```
==17647==    by 0x4006C0: main (foo.c:15)
```

```
==17647== Address 0x5200045 is 0 bytes after a block of size 5 alloc'd
```

```
==17647==    at 0x4C2DB9D: malloc (vg_replace_malloc.c:299)
```

```
==17647==    by 0x400645: string_copy (foo.c:7)
```

```
==17647==    by 0x400690: main (foo.c:13)
```

```
==17647==
```

```
hello
```

```
==17647==
```

*first half of output*

# Valgrind

```
1. #include <stdio.h>
2. #include <stdlib.h>
3. #include <string.h>
4. #include <assert.h>
5. char *string_copy( const char *in )
6. {
7.     char *out = malloc( strlen( in ) );
8.     assert( out!= NULL );
9.     return strcpy( out , in );
10. }
11. int main( void )
12. {
13.     char *str = string_copy( "hello" );
14.     assert( str!=NULL );
15.     printf( "%s\n" , str );
16.     return 0;
17. }
```

```
>> ./a.out
hello
>>
```

*second half of output*

```
==17647== HEAP SUMMARY:
==17647==     in use at exit: 5 bytes in 1 blocks
==17647==   total heap usage: 2 allocs, 1 frees, 1,029 bytes allocated
==17647==
==17647== 5 bytes in 1 blocks are definitely lost in loss record 1 of 1
==17647==    at 0x4C2DB9D: malloc (vg_replace_malloc.c:299)
==17647==    by 0x400645: string_copy (foo.c:7)
==17647==    by 0x400690: main (foo.c:13)
==17647==
==17647== LEAK SUMMARY:
==17647==    definitely lost: 5 bytes in 1 blocks
==17647==    indirectly lost: 0 bytes in 0 blocks
==17647==    possibly lost: 0 bytes in 0 blocks
==17647==    still reachable: 0 bytes in 0 blocks
==17647==         suppressed: 0 bytes in 0 blocks
==17647==
==17647== For counts of detected and suppressed errors, rerun with: -v
==17647== ERROR SUMMARY: 3 errors from 3 contexts (suppressed: 0 from 0)
>>
```

# Valgrind

- So what was wrong?
  - An invalid write
  - An invalid read
  - A block of memory that wasn't *freed*
- On `ugradx`, this program didn't crash and seemed to work properly
  - ⇒ Not every bad memory access leads to error (or bad output)
  - ⇒ `valgrind` is a useful tool to help us find problematic code

# Valgrind

```
1. #include <stdio.h>
2. #include <stdlib.h>
3. #include <string.h>
4. #include <assert.h>
5. char *string_copy( const char *in )
6. {
7.     char *out = malloc( strlen( in )+1 );
8.     assert( out!= NULL );
9.     return strcpy( out , in );
10. }
11. int main( void )
12. {
13.     char *str = string_copy( "hello" );
14.     assert( str!=NULL );
15.     printf( "%s\n" , str );
16.     free( str );
17.     return 0;
18. }
```

```
>> ./a.out
hello
>>
```

```
>> valgrind --leak-check=full ./a.out
                                     header
==30398== Command: ./a.out
==30398==
hello
==30398==
==30398== HEAP SUMMARY:
==30398==     in use at exit: 0 bytes in 0 blocks
==30398==   total heap usage: 2 allocs, 2 frees, 1,030 bytes allocated
==30398==
==30398== All heap blocks were freed -- no leaks are possible
==30398==
==30398== For counts of detected and suppressed errors, rerun with: -v
==30398== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
>>
```



# Outline

- Dynamic memory allocation
- Valgrind
- Review questions

# Review questions

1. What is the difference between stack and heap memory?

# Review questions

2. What is dynamic memory allocation in C?

# Review questions

3. What is the memory leak problem?

# Review questions

4. What is the difference between `malloc`, `realloc`, and `calloc`?

# Review questions

5. What do we use `valgrind` to check for?

# Review questions

6. Consider the `exclaim` function below. Do you see any problems with this function?

```
// Return a C character string containing n exclamation points.  
// n must be less than 20.  
char *exclaim( int n )  
{  
    char s[20];  
    assert( n<20 );  
    for( int i=0 ; i<n ; i++ ) s[i] = '!';  
    s[n] = '\\0';  
    return s;  
}
```

# Exercise 4-2

- Website -> Course Materials -> Ex4-2