600.220 Intermediate Programming

Type conversions

Plan for today

- Random integers in C
- Number type representations
- Casting between types

Random integers in C

- rand() generates (pseudo) random integers between 0 and RAND_MAX
- the pseudo sequence of random integers is based on a time seed
- srand(unsigned int) sets the time seed
- if srand() is not called, by default, it uses time seed 1 (as if called srand(1) at the beginning of the program)
- use srand(time(0)) to generate time dependent random integeters (time.h is required)

Random integers in C

```c
// random.c:
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

int main() {
    for (int i = 0; i < 5; ++i)
        printf(" %d ", rand()); // print 5 random integers w/o calling srand()
    printf("\n");
    srand(time(0)); // Set seed to current time
    for (int i = 0; i < 5; ++i)
        printf(" %d ", rand()); // print another 5 random integers
    printf("\n");
    srand(1); // Set seed back to 1
    for (int i = 0; i < 5; ++i)
        printf(" %d ", rand()); // print another 5 random integers
    printf("\n");
    return 0;
}
```

```
$ gcc -std=c99 -pedantic -Wall -Wextra -c random.c
$ gcc -o random random.o
$ ./random
 1804289383  846930886  1681692777  1714636915  1957747793
 256385480  614463771  393795659  220264062  212663161
 1804289383  846930886  1681692777  1714636915  1957747793
```

Number representations: integers

C integers use "two's complement" representation for signed
integers. Illustration with 4 bits:

http://www.bogotobogo.com/cplusplus/quiz_bit_manipulation.php

Integer overflow

When a two's complement number overflows, it wraps around to a negative number

```
// overflow.c:
#include <stdio.h>

int main() {
    int i = 2147483647;  // largest integer value
    int i_plus_1 = i + 1;
    printf("i = %d, i+1 = %d\n", i, i_plus_1);
    return 0;
}

$ gcc -std=c99 -pedantic -Wall -Wextra -c overflow.c
$ gcc -o overflow overflow.o
$ ./overflow
i = 2147483647, i+1 = -2147483648
```

Number representations: floating-point numbers

Floating point numbers use their bits to store a few different things, e.g. (IEEE-754 standard):

- Sign: 1 bit, positive or negative
- Exponent: 8 bits (float), 11 bits (double)
- Mantissa: 23 bits (float), 52 bits (double)

sign exponent mantissa

| 0 | 1 | 1 | 0 | 1 | 0 | 1 | 1 |

$\downarrow$

$+ \ (0.1011)_2 * 10^{(110)_{\text{3bit excess-k}}}$

$= (0.1011)_2 * 10^2$

$= (10.11)_2$

$= (1*2)^1 + (1*2)^0 + (0*2)^{-1} + (1*2)^{-2} = 2.75$

https://aha.betterexplained.com/t/8bit-floating-point-representation/1121

Number representation differences

Integer and floating-point representations differ:

- Integers have limited range, but integers in the range can be represented precisely. Floating point have limited range and can only approximate most numbers in the range.
- Integers use all available bits for two's-complement representation. Floating point have separate sets of bits for sign, exponent and mantissa.

Conversions between types

- When you do `float a = 1` or `int i = 3.0`, it's not as simple as copying bits

- Conversion from an integer type to `float` or `double` might not be exact (`float` can exactly represent integers in the range $\pm 2^{24}$, `double` can exactly represent integers in the range $\pm 2^{53}$)

- Our CSF (Computer System Fundamentals) course goes into detail on conversions and arithmetic involving these representations

Integer literal types

- The type for an *integer literal* (e.g. 88, -1000000000) is determined based on its value
    - Specifically: its type is the smallest integer type (starting from int) that can store it without overflowing

Integer literal types

```c
// int_literal.c:
#include <stdio.h>

int main() {
    int a = 1;
    int b = -3000;
    long c = 10000000000; // too big for an int
    printf("%d, %d, %ld\n", a, b, c);
    return 0;
}

$ gcc -std=c99 -pedantic -Wall -Wextra -c int_literal.c
$ gcc -o int_literal int_literal.o
$ ./int_literal
1, -3000, 10000000000
```

Integer literal range

Compiler can warn us when an integer literal is too big for a particular variable

```c
// int_literal2.c:
#include <stdio.h>

int main() {
    int a = 10000000000;  // too big for an int
    // a will "overflow" and "wrap around" to some other number
    printf("%d\n", a);
    return 0;
}

$ gcc -std=c99 -pedantic -Wall -Wextra -c int_literal2.c
int_literal2.c: In function 'main':
int_literal2.c:4:13: warning: overflow in implicit constant conversion [-Woverf
    int a = 10000000000;  // too big for an int
            ^
$ gcc -o int_literal2 int_literal2.o
$ ./int_literal2
1410065408
```

Floating-point literal types

Floating-point literal (e.g. 3.14, −.7, −1.1e−12) has type double
You can force it to be float by adding f suffix

```c
// float_literal.c:
#include <stdio.h>

int main() {
    float a = 3.14f;
    double b = 33.33, c = -1.1e-12;
    printf("%f, %f, %e\n", a, b, c);
    return 0;
}
```

```
$ gcc -std=c99 -pedantic -Wall -Wextra -c float_literal.c
$ gcc -o float_literal float_literal.o
$ ./float_literal
3.140000, 33.330000, -1.100000e-12
```

Floating-point literal range

Again, compiler will warn if literal doesn't fit

```c
// float_literal2.c:
#include <stdio.h>

int main() {
    float a =  2e128f; // too big; max float exponent is 127
    double b = 2e1024; // too big; max double exponent is 1023
    printf("%f, %f\n", a, b);
    return 0;
}
```

Floating-point literal range

```
$ gcc -std=c99 -pedantic -Wall -Wextra -c float_literal2.c
float_literal2.c: In function 'main':
float_literal2.c:4:5: warning: floating constant exceeds range of 'floa
     float a =  2e128f; // too big; max float exponent is 127
     ^
float_literal2.c:5:5: warning: floating constant exceeds range of 'doub
     double b = 2e1024; // too big; max double exponent is 1023
     ^
$ gcc -o float_literal2 float_literal2.o
$ ./float_literal2
inf, inf
```

Zoom poll!

Assume char is 8 bits and is a signed type, sizeof(short)=2, sizeof(int)=4, and sizeof(long)=8. Also assume two's complement representation. Consider the following code:

```
char p = 128;
short q = -30197;   // -30,197
int r = 3000000000; // 3,000,000,000
long s = (1L << 33) * (1L << 33);
```

Which variables are assigned values that exceed their range?

A. p
B. p and r
C. q and r
D. p, r, and s
E. q, r, and s

Type Conversions

C can automatically convert between types "behind the scenes"
This is called *automatic conversion* and can be a *promotion* in the
case of a smaller type value converted to a larger type, or *narrowing*
in the case of a larger type value converted to a smaller type.

```
float ten = 10;
// float <- int

int ten = 10.585; // truncates to 10
// int <- float
```

Promotion

```c
// promotion_1.c:
#include <stdio.h>

int main() {
    int a = 1;
    float f = a * 1.5f;
    printf("%f\n", f);
    return 0;
}

$ gcc -std=c99 -pedantic -Wall -Wextra -c promotion_1.c
$ gcc -o promotion_1 promotion_1.o
$ ./promotion_1
1.500000
```

Promotion

```
int a = 1;
float f = a * 1.5f;
```

Note operands types: a is int, 1.5f is float

When operand types don't match, "smaller" type is promoted to "larger" before applying operator

char < int < unsigned < long < float < double

Promotion

```c
// promotion_2.c:
#include <stdio.h>

int main() {
    int a = 3;
    float f = a / 2;
    printf("%f\n", f);
    return 0;
}

$ gcc -std=c99 -pedantic -Wall -Wextra -c promotion_2.c
$ gcc -o promotion_2 promotion_2.o
$ ./promotion_2
1.000000
```

Promotion

```
int a = 3;
float f = a / 2;
```

No promotion here before division, since operands a and 2 are same
type: int

Narrowing

Type conversions from larger to smaller types (e.g. double -> float or long -> int) can happen automatically too - explicit type casts are not required!

Sometimes called *narrowing* conversions, these usually chop off the extra bits without rounding values.

```c
// narrow_1.c:
#include <stdio.h>

int main() {
    unsigned long a = 1000;
    int b = a; // automatic *narrowing* conversion
    float c = 3.14f;
    double d = c; // automatic conversion
    printf("b=%d, d=%f\n", b, d);
    return 0;
}
```

Narrowing

```
$ gcc -std=c99 -pedantic -Wall -Wextra -c narrow_1.c
$ gcc -o narrow_1 narrow_1.o
$ ./narrow_1
b=1000, d=3.140000
```

No compiler warnings.

Narrowing

```
// narrow_2.c:
#include <stdio.h>
int square(int num) {
    return num * num;
}
int main() {
    printf("square(2.5)=%d\n", square(2.5));
    return 0;
}
```

Narrowing

```
$ gcc -std=c99 -pedantic -Wall -Wextra -c narrow_2.c
$ gcc -o narrow_2 narrow_2.o
$ ./narrow_2
square(2.5)=4
```

2.5 becomes 2 when passed to square. Again no compiler warning.

Narrowing

A value's type is narrowed *automatically* and *without a compiler warning* when: (a) assigning to a variable of narrower type, and (b) passing an argument into a parameter of narrower type.

Other automatic narrowing situations typically yield compiler warnings - you should eliminate the warnings in your programs by using explicit type casts!

Narrowing

```c
// casting_1.c:
#include <stdio.h>

int main() {
    printf("sizeof(long)=%d\n", sizeof(long));
    return 0;
}
```

```
$ gcc -std=c99 -pedantic -Wall -Wextra -c casting_1.c
casting_1.c: In function 'main':
casting_1.c:4:12: warning: format '%d' expects argument of type 'int',
    printf("sizeof(long)=%d\n", sizeof(long));
            ^
$ gcc -o casting_1 casting_1.o
$ ./casting_1
sizeof(long)=8
```

Casting

Some types just can't be used for certain things. E.g. a `float` can't be an array index:

```c
// casting_2.c:
#include <stdio.h>

int main() {
    int array[] = {2, 4, 6, 8};
    float f = 3.0f;
    printf("array[0]=%d, array[%f]=%d\n", array[0], f, array[f]);
    return 0;
}
```

Casting

```
$ gcc -std=c99 -pedantic -Wall -Wextra -c casting_2.c
casting_2.c: In function 'main':
casting_2.c:6:61: error: array subscript is not an integer
     printf("array[0]=%d, array[%f]=%d\n", array[0], f, array[f]);
                                                               ^
```

Casting

Type *casting* gives you more control over when promotion and narrowing happens in your program

Casting is sometimes the only way to avoid compiler errors and warnings

Even when conversion would happen automatically, making it explicit by using a cast can make your code clearer

Casting is a higher precedence operation than the binary arithmetic operators since it is a unary operation

Casting

```c
// casting_3.c:
#include <stdio.h>

int main() {
    int a = 3;
    float f = (float) a / 2;
    //        ^^^^^^^
    // a gets *cast* to float
    // 2 gets *promoted* to float before division
    printf("%f\n", f);
    return 0;
}

$ gcc -std=c99 -pedantic -Wall -Wextra -c casting_3.c
$ gcc -o casting_3 casting_3.o
$ ./casting_3
1.500000
```

Casting

```
// casting_4.c:
#include <stdio.h>

int main() {
    printf("sizeof(long)=%d\n", (int)sizeof(long));
    //                          ^^^^^
    return 0;
}

$ gcc -std=c99 -pedantic -Wall -Wextra -c casting_4.c
$ gcc -o casting_4 casting_4.o
$ ./casting_4
sizeof(long)=8
```

Casting

```
// casting_5.c:
#include <stdio.h>

int main() {
    int array[] = {2, 4, 6, 8};
    float f = 3.0f;
    printf("array[0]=%d, array[%d]=%d\n",
           array[0], (int)f, array[(int)f]);
    //
    return 0;
}

$ gcc -std=c99 -pedantic -Wall -Wextra -c casting_5.c
$ gcc -o casting_5 casting_5.o -lm
$ ./casting_5
array[0]=2, array[3]=8
```

Zoom poll!

What output is printed by the following program?

```c
#include <stdio.h>

int main(void) {
  int x = 20000;
  // 20,000 decimal is
  //   01001110 00100000
  // in binary
  printf("%d\n", (unsigned char) x);
  return 0;
}
```

A. 0
B. 16
C. 20
D. 32
E. 20000

Type mystery

```
// casting_6.c:
#include <stdio.h>
#include <math.h>

int main() {
    float p = 2000.0, r = 0.10;
    float ci_1 = p * pow(1 + r, 10);
    float ci_2 = p * pow(1.0f + r, 10);
    float ci_3 = p * pow(1.0 + r, 10);
    printf("%.3f\n%.3f\n%.3f\n", ci_1, ci_2, ci_3);
    return 0;
}

$ gcc -std=c99 -pedantic -Wall -Wextra -c casting_6.c
$ gcc -o casting_6 casting_6.o -lm
$ ./casting_6
5187.486
5187.486
5187.485
```

Type mystery: solved

Prototype for pow: `double pow(double, double);`
Type promotions are happening both because of the addition and because of the call to pow

- `ci_1`: 1 converted to `float`, then added to `r`, then result is converted to `double`
- `ci_2`: `1.0f + r` converted to `double`
- `ci_3`: `r` converted to `double`, then added to `1.0` (already a `double`)

`(float)1` and `(double)1` are not the same. `ci_1` and `ci_2` use `(float)1`, `ci_3` uses `(double)1`.